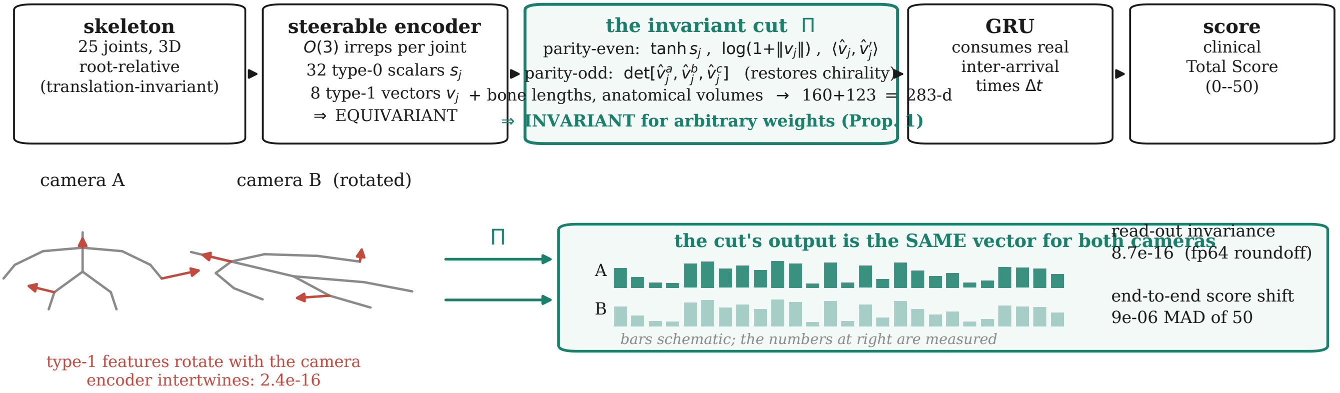
skeleton
25 joints, 3D
root-relative
(translation-invariant)
steerable encoder
$O(3)$ irreps per joint
32 type-0 scalars s_j
8 type-1 vectors v_j
⇒ EQUIVARIANT
the invariant cut Π
parity-even: $\tanh s_j$, $\log(1+\|v_j\|)$, $\langle \hat{v}_j, \hat{v}'_j \rangle$
parity-odd: $\det[\hat{v}^a_j, \hat{v}^b_j, \hat{v}^c_j]$ (restores chirality)
+ bone lengths, anatomical volumes → 160+123 = 283-d
⇒ INVARIANT for arbitrary weights (Prop. 1)
GRU
consumes real
inter-arrival
times Δt
score
clinical
Total Score
(0--50)
camera A
camera B (rotated)
type-1 features rotate with the camera
encoder intertwines: 2.4e-16
Π
the cut's output is the SAME vector for both cameras
A
B
bars schematic; the numbers at right are measured
read-out invariance
8.7e-16 (fp64 roundoff)
end-to-end score shift
9e-06 MAD of 50