

Viewpoint Invariance Is a Theorem, Not a Fit: An $SE(3)$ -Equivariant Recurrence for Rehabilitation Exercise Assessment

Anonymous WACV Applications Track submission

Paper ID 0000

Abstract

A clinical screening instrument must return the same patient the same score from a different room. We show that for skeleton-based rehabilitation assessment this per-prediction guarantee is available by construction: an $SE(3)$ -equivariant recurrent model whose read-out is invariant to camera viewpoint as a theorem, certified to the fp64 roundoff floor ($\sim 10^{-16}$) over Haar-random $SO(3)$ — of which any camera relocation, including the 45° NTU RGB+D Cross-View displacement, is one instance. On clean frontal data no model separates from another — ours, a point-cloud-transformer baseline, and a hand-crafted-invariant GRU tie within a 0.33 MAD nondeterminism floor — so the benchmark’s aggregate metric is saturated and the decisive axis lies elsewhere: the augmented baseline matches us in aggregate accuracy while still swinging an individual patient’s score by up to 20.1 MAD (of 50) under camera rotation, against our 9×10^{-6} . We then map the accuracy–exactness–robustness–cost frontier and find that no model on it dominates: the point-cloud baseline is both the most accurate and the most node-robust, a hand-crafted-invariant control is the cheapest and exactly invariant yet collapses under sensor-node failure, and a lighter $E(n)$ encoder behind the identical invariant cut is measurably more brittle. Ours is the only point that is exact by theorem while clearing every stress, at $7.4\times$ fewer parameters. We replicate the viewpoint and node-failure properties on a second corpus (REHAB24-6), and measure what the theorem costs on the edge: an int8 activation-quantization cliff of only 0.05 MAD — 0.005 of a between-subject SD on this corpus, and one to two orders of magnitude below MCID thresholds on comparable rehabilitation scales (Sec. 4) — and a causal streaming variant at 112 fps (GPU, batch 1). We report our negative results — a null irregular-sampling regime and an unrewarded parity restoration — with equal prominence.

1. Introduction

Automated assessment of rehabilitation exercises from skeletons promises objective remote monitoring [1]. The dominant regime, however, trains and tests on clean frontal recordings from one sensor and ranks models by a single aggregate error against a clinician’s score [2–5]. It rewards capacity and penalizes nothing a deployed device meets: a camera against a different wall, a tracking node that freezes, a dropped frame.

This paper argues that the *regime*, not the architectures, is the problem. On clean frontal KIMORE no architecture separates from another: our model, the point-cloud transformer, and a GRU on six hand-crafted invariant features are statistically indistinguishable — the benchmark’s aggregate metric is saturated. We can say so precisely because we first measured what a difference must exceed to be one: training here is not bit-reproducible (nondeterministic cuDNN and scatter-add kernels), and repeated fits of an identical configuration differ by ≈ 0.33 MAD — larger than every clean-accuracy gap between the leading models. A single-run comparison on this benchmark reports its own seed. Worse, auditing the reference transformer on the commonly used single-exercise slice, optimistic epoch selection is worth $+1.2 \pm 0.4$ MAD (18.6% over three seeds), and its honest number carries no reliable subject-level signal: the paired subject-cluster bootstrap gives $\Delta = -0.44$ (per-seed $[-1.07, +0.20]$), clearing zero in only one of three seeds — the slice’s limit, not the model’s. We therefore pool the five exercises, keep splits subject-disjoint, and gate every claim on clearing a per-exercise mean-predictor floor — a model that ignores its input has zero degradation and would win every robustness plot. Pooled, every model we test clears it (Tab. 3).

With clean accuracy neutralized, the question becomes what a structural guarantee buys. We formulate assessment as an invariant read-out of an equivariant representation: a steerable message-passing encoder produces per-joint features in $O(3)$ irreducible representations, an explicit pro-

jection onto a generating set of invariants discharges the symmetry, and a recurrence consumes the invariant code with the sensor’s actual inter-arrival times (a capability KIMORE does not reward, Sec. 4; kept because it is free and demonstrably real off this corpus). Invariance of the score under any rigid camera motion is then a *theorem* (Proposition 1), true for arbitrary weights, not a behaviour induced by augmentation.

We contribute three things, none of them new invariance theory (scalar contractions of Wigner- D -covariant features are invariant by construction [10–12]). An *end-to-end certification* protocol: eight numerical gates that survive the real steerable layers, the adaptive solver and int8 quantization. An *evaluation methodology* the field lacks: a mean-predictor floor below which no gap is interpretable, and per-sequence degradation reported as a distribution rather than a mean. And a *deployment* analysis: a precision budget and a causal streaming read-out. The distinction bites: a rotation-augmented baseline has a flat *aggregate* accuracy curve across azimuth — by the field’s metric it has solved viewpoint — yet its *per-sequence* degradation averages 3.0 MAD and reaches 20.1 in the tail, individual scores swinging up to two fifths of the fifty-point scale as the camera moves, the errors merely cancelling in the mean. Ours move by 9×10^{-6} . We therefore report degradation and accuracy side by side throughout.

The tie with the hand-crafted model leaves an obligation elegance cannot discharge: if a plain GRU on named features is exactly as accurate *and* as viewpoint-invariant, the steerable machinery is unjustified. We discharge it against *that* model specifically. When Kinect nodes freeze — the characteristic consumer-depth failure — the hand-crafted model fails outright: a *single* dead joint drives it from 6.31 to 9.78 MAD, through the floor, because its features *name* the joints they depend on; the graph encoder pools learned messages and dilutes the failure across 25 joints. That is one pairwise comparison, not an axis we own: the point-cloud baseline is more node-robust than either of us, and our case rests on *where* we sit on the accuracy–exactness–robustness–cost frontier (Tab. 6). We also treat the guarantee as a *systems* property: it survives to int8 (only activation, not weight-rounding perturbs it) and converts to strictly causal streaming at 8.9 ms/frame. Finally we report, in full, the two places the structural story does not pay — native consumption of irregular arrivals, which a fixed grid beats at every drop level here, and a restored parity channel that is principled but unrewarded on healthy-form exercises.

Contributions. (i) A protocol audit that establishes what the benchmark can measure — an 18.6% epoch-selection inflation on the reference model, a 0.33 MAD nondeterminism floor below which no accuracy gap is interpretable, and a gated pooled protocol under which every model clears the floor. (ii) Exact viewpoint invariance, *certified* by eight

numerical gates (E1–E8) over Haar-random $SO(3)$ rather than a single azimuth sweep, against a baseline that crosses the floor by 45° . (iii) A Pareto characterisation of the accuracy–exactness–robustness–cost frontier on which no model dominates: the baseline wins clean accuracy *and* node robustness, while ours is alone in being exact by theorem at $7.4\times$ fewer parameters — including, behind the identical invariant cut, evidence that lighter $E(n)$ equivariance is more node-fail brittle. (iv) A second-corpus replication (REHAB24-6) of the viewpoint and node-failure properties. (v) A co-designed deployment path: an int8 precision-budget audit and a causal streaming read-out with an honest time-to-first-score. We also report two negative results, one of them derived.

2. Related Work

Rehabilitation assessment and equivariant models. KIMORE [1] is the standard benchmark; reported results span ST-GCN variants with contrastive objectives [3], cross-modal augmentation [4], dual-stream graph networks [5], and point-cloud transformers [2], with Rehab-Pile [6] (aggregating KIMORE, UI-PRMD [7] and IRDS) the closest antecedent of our protocol discipline. Evaluation is uniformly in-distribution (same sensor, same placement, no simulated failure), and to our knowledge no prior work reports a mean-predictor floor, without which a degradation curve cannot be interpreted.

Published numbers do not enter Tab. 3, because no protocol matches ours. Kuang et al. [5] report MAD 0.10–0.16 in the *same* 0–50 clinical-score units as our subject-disjoint 6.3–6.7, lower by a factor of 39–67; their split is an explicit 8:1:1 *sample-level* stratified one. We recover those units from the published tables (supplement) because the reference implementation standardises its targets and reading it wrong would move the comparison by $8.5\times$. Karlov et al. [3] report Spearman ρ rather than MAD, so we report ρ too: our pooled $\rho \approx 0.6$ sits below the 0.74–0.965 band the literature spans, on a protocol whose mean-predictor null we also state. Abedi et al. [4] report only a relative MAD reduction with no recoverable absolute baseline. Lacking their preprocessed data and a subject-disjoint rerun, we treat these as measured under a different protocol, not as an accuracy ceiling.

Group-equivariant networks impose symmetry as a constraint on the hypothesis class [9]; the steerable family — tensor-field networks [10], $SE(3)$ -Transformers [11], e3nn [12] — guarantees equivariance for arbitrary weights, while lighter constructions reach $E(n)$ -equivariance without spherical harmonics [13] or via vector neurons [14]. We build on [12], but the claim is *not* about steerable machinery: a GRU on hand-crafted invariants attains the same accuracy and viewpoint guarantee at a third of the parameters, and we isolate the one axis on which they differ (Tab. 4).

Invariance by canonicalization, augmentation, and corruption robustness. A non-equivariant backbone can be made invariant by canonicalization or frame averaging [15], exact only in so far as the frame estimate is — a point estimate whose *conditioning* degrades near covariance degeneracy, which we measure directly (Sec. 4); our cut estimates no frame, being an explicit projection exact for arbitrary weights (Proposition 1). Training-time augmentation flattens *aggregate* error but makes per-prediction errors cancel, not vanish — expressible only if degradation is reported per sequence. For irregular sampling, the standard alternatives are resampling or continuous-time models [18–21]; we implement the Neural CDE variant (adaptive Dormand–Prince solver [22]) as a certified *control* that fails the floor. Occlusion and tracking failure are usually met with masking or attention [8]; we adopt the failure literally and use it as the discriminating experiment between two otherwise-indistinguishable models.

3. Method

A steerable encoder on the skeleton graph carries per-joint latents in $O(3)$ irreps (n_0 type-0 scalars s_j , n_1 type-1 vectors v_j); feeding *root-relative* coordinates makes translation invariance automatic. The encoder is *equivariant*; the score must be *invariant*.

The invariant cut. The transition is an explicit projection Π onto a generating set of invariants. The parity-even generators (scalars, vector norms, pairwise cosines) are complete for the full $O(3)$ — each is fixed by a reflection — so a cut built from them is blind to *trajectory handedness* (a clockwise and anticlockwise limb circle share every bone length, distance and speed). We complete it with parity-odd pseudo-scalar triple products over unit vectors:

$$\pi_j^{\text{even}} = \left[\tanh s_j \parallel \log(1 + \|v_j^c\|) \parallel \langle \hat{v}_j^c, \hat{v}_{j'}^c \rangle_{c < c'} \right] \quad (1)$$

$$\pi_j^{\text{odd}} = \left[\det[\hat{v}_j^a, \hat{v}_j^b, \hat{v}_j^c]_{a < b < c} \right] \quad (2)$$

$$\pi_j = [\pi_j^{\text{even}} \parallel \pi_j^{\text{odd}}] \quad (3)$$

with $\hat{v} = v/\|v\|$, pooled over joints by $\text{mean} \oplus \text{max}$ and concatenated with invariant bone lengths and a fixed set of anatomical signed volumes χ_τ (index-preserved, unpooled):

$$\Pi(h, \tilde{x}) = \left[\text{mean}_j \pi_j \parallel \max_j \pi_j \parallel (d_{ij})_\mathcal{E} \parallel (\chi_\tau)_\mathcal{T} \right]. \quad (4)$$

At the deployed widths ($n_0=32$, $n_1=8$) the parity-even subtotal is 160-d and the parity-odd completion adds 123-d (283-d total; full specification in the supplement). A liveness gate zeroes any generator touching a failed joint, which makes the dead-node *invariance* certified by E1–E8 exact under a known failure mask; the degradation reported in

Table 1. **The viewpoint guarantee is falsifiable.** Eight numerical gates on the deployed read-out: E1–E5 that the invariance is *present*, E6–E8 that it is the *intended* one (chirality live, laterality preserved). Roundoff-floor quantities are regime-stable, not bit-reproducible (representative fp64 run); E6–E7 are *modelling* quantities, not roundoff, reported for seed 0 of `certify_egr_u.json`. Parity/laterality gates report the even/chiral cut.

Gate	Criterion	Result
E1	rep. orthogonality $\ \rho^\top \rho - I\ _F$	$\leq 10^{-12}$ ✓
E2	motion-channel invariance	verified ✓
E3	rotation-sweep violation / slope	$\sim 10^{-16}$, $ \cdot < 0.5$ ✓
E4	fp32/fp64 violation ratio r	$> 10^8$ ✓
E5	translation invariance	$\sim 10^{-16}$ ✓
E6	parity (even / chiral)	$0 / 4.4 \times 10^{-2}$
E7	laterality (even / chiral)	$6.9 / 5.6 \times 10^{-2}$
E8	full- $SO(3)$, parity-odd active	$\sim 10^{-16}$ ✓

Tab. 4 instead uses the deployed gate-*off* model with no failure signal supplied, so the graceful fall-off measured there is message-passing dilution alone, not the gate.

Proposition 1 ($SO(3)$ -invariance of the read-out) *For an encoder that intertwines with the Wigner-D action on its latents, Π of Eq. (3)–(4) is invariant under $SO(3)$ for arbitrary weights; the predicted score is therefore invariant to any rigid motion of the camera. (Proof in the supplement.)*

Certificate. We audit the read-out with eight numerical gates (Tab. 1), all passed, each independently re-runnable and refutable (see Code availability): representation orthogonality $\leq 10^{-12}$; violation at the fp64 roundoff floor ($\sim 10^{-16}$) with a *flat* log-log slope in rotation angle ($|\text{slope}| < 0.5$, against ≈ 1 for a genuine modelling error) and an fp32/fp64 violation ratio $> 10^8$ (excluding a small modelling violation masquerading as noise); translation and read-out invariance both $\sim 10^{-16}$ (fp64) after the parity-odd channels are admitted. The guarantee is over Haar-random proper rotations of the *full* $SO(3)$, of which a camera azimuth or a cross-view displacement is one instance. A continuous-time (Neural CDE) variant, certified to the same standard, *fails* the mean-predictor floor and is reported only as a control (supplement).

4. Results

All experiments are pooled, subject-disjoint, over 3 seeds, except the TCN, ST-GCN and Ridge viewpoint curves of Fig. 2, which are 5-fold cross-validated at a single seed. *Protocol audit* (Tab. 2): we audit the point-cloud transformer baseline on the single-exercise slice most commonly reported. Honest epoch selection removes an 18.6% inflation (3 seeds) and leaves that model with no robust separation from a mean predictor — the paired subject-cluster




Figure 1. **The invariant cut is the method.** Top: the pipeline. Bottom: the same pose at two camera poses — the encoder’s type-1 features rotate (equivariance), the cut’s output does not (invariance, Prop. 1). Bar heights are schematic; the annotated quantities are measured (certify_egru.json, final_tables.json).

Table 2. Protocol audit of the *point-cloud transformer baseline*, single-exercise slice (KIMORE Ex. 1; subject-disjoint 5-fold; 3 seeds). *Honest* selects the epoch on a held-out validation fold; *inflated* takes the per-epoch test minimum (the reference protocol). Δ is the paired subject-cluster bootstrap MAD(honest) – MAD(floor); an interval below 0 means the model beats the mean predictor. The signal here is *seed-fragile*, clearing zero in only 1/3 seeds; pooling fixes this, for this model and every other (Tab. 3).

Seed	Honest / Inflated	Selection bias	Δ vs. floor (95% CI)
0	6.90/5.45	+1.44 (20.9%)	+0.20 [−0.59, +1.03]
1	6.53/5.00	+1.54 (23.5%)	−0.46 [−1.60, +0.72]
2	5.84/5.17	+0.67 (11.4%)	−1.07 [−1.98, −0.15]
mean	6.42±0.44/5.21±0.19	+1.22 (18.6%)	−0.44

bootstrap gives $\Delta = -0.44$ and clears zero in only 1/3 seeds. This is not subject leakage (each subject contributes one recording on that slice, so a random split is already subject-disjoint); the slice simply carries no robust subject-level signal, which is why we pool the five exercises and gate every experiment on a per-exercise mean-predictor floor. Pooled, every model here clears that floor (Tab. 3).

The clean-accuracy metric is saturated. EGRU (6.73 ± 0.30), a GRU on hand-crafted invariants (Invariant-GRU, 6.31 ± 0.16), and the point-cloud transformer (PCT, 6.47 ± 0.20) are mutually indistinguishable: every gap is inside the 0.33 MAD seed floor (Tab. 3). This is a statement about the models *relative to each other*, not about whether any of them carries real subject-level signal: unlike the fragile single-exercise slice (Tab. 2, clears zero in only 1/3 seeds), the pooled bootstrap Δ vs. floor is entirely below 0 for all three — PCT $-1.87 [-2.72, -1.06]$, Invariant-GRU $-2.03 [-2.88, -1.23]$, EGRU $-1.60 [-2.37, -0.88]$ — so the tie is a tie among models that each genuinely beat the mean predictor, not three models indistinguishable from noise. The tie survives the metric change: Spearman ρ separates the same models by at most 0.05 (Tab. 3), so saturation is a property of the benchmark, not of the error measure

chosen to report it. ρ also needs its null stated — the mean predictor scores pooled $\rho = 0.17$ while ignoring its input entirely, and exactly 0 once the between-exercise score differences it exploits are removed. A benchmark on which three architectures of vastly different inductive bias cannot be told apart is not measuring architecture — which is the premise for evaluating on an axis that can. Restoring chirality ($O(3) \rightarrow SO(3)$) costs $+0.11$ MAD for EGRU — inside the floor — for 16.7% more parameters and the removal of a silent handedness assumption; we adopt the chiral model on that basis, not an accuracy one. The chiral read-out’s wider seed spread (± 0.30 vs ± 0.08 for the parity-even model) reflects the added parity-odd capacity, not instability — every seed stays inside the floor.

One asymmetry against the baseline is worth closing: EGRU sees a one-hot exercise id and the floor is per-exercise, so pooled PCT alone was scored blind to the exercise. Re-training it with identical conditioning (3 seeds, both arms) changes nothing: clean 6.59 ± 0.07 vs. 6.47 ± 0.20 , augmented 7.42 ± 0.22 vs. 7.47 ± 0.23 — inside the floor and opposite in sign — and the per-sequence swing is if anything *larger* (21.43 vs. 20.12 MAD worst sequence). Viewpoint fragility is not an artefact of withholding exercise identity; we report the unconditioned arm throughout.

Viewpoint invariance is a theorem, not a fit (Fig. 2). Across test-time camera azimuth the certified-invariant models are flat to machine precision (EGRU mean 9×10^{-6} MAD, worst single sequence 3×10^{-4} ; Invariant-GRU exactly 0; hand-crafted invariant features Ridge at the fp64 roundoff floor, 10^{-14} — both invariant by construction, so these two curves are a positive control, not independent evidence), while all non-equivariant baselines degrade catastrophically: PCT crosses the mean-predictor floor between 30° and 45° (mean per-sequence shift 9.33 MAD at its worst angle, worst sequence 33.96 of fifty); TCN 13.17 MAD; ST-GCN 9.18 MAD. This pattern holds across

Table 3. Clean-data accuracy (pooled, subject-disjoint, mean $\pm$ std over 3 seeds). The 0.33 MAD floor is the *fixed-configuration* one; the wider seed-to-seed spread (0.48) is in the supplement. Δ vs. floor is the paired subject-cluster bootstrap of Tab. 2 ($n=77$ subjects, checkpoints re-evaluated without retraining; `pooled_bootstrap_eval.py`). Spearman ρ is computed per seed on that seed’s complete out-of-fold vector, pooled over exercises and again within each exercise.

Model	Group	MAD $\downarrow$	Params	Spearman $\rho \uparrow$		Δ vs. floor (95% CI)
				pooled	within-ex.	
PCT (baseline)	—	6.47 $\pm$ 0.20	4.91M	0.60 $\pm$ 0.02	0.55 $\pm$ 0.02	-1.87 [-2.72, -1.06]
PCT + rotation aug.	—	7.47 $\pm$ 0.23	4.91M	—	—	—
InvariantGRU (hand-crafted)	SO(3)	6.31 $\pm$ 0.16	0.21M	0.59 $\pm$ 0.01	0.56 $\pm$ 0.01	-2.03 [-2.89, -1.23]
EGRU (ours)	SO(3)	6.73 $\pm$ 0.30	0.66M	0.56 $\pm$ 0.04	0.52 $\pm$ 0.04	-1.60 [-2.37, -0.88]
Mean-predictor floor	—	8.31 $\pm$ 0.05	—	0.17 $\pm$ 0.01	-0.07 $\pm$ 0.02	—

three independently-designed non-equivariant architectures (temporal-conv, graph-conv, attention), indicating viewpoint fragility is a generic property of non-equivariance, not PCT-specific. Rotation augmentation flattens PCT’s aggregate curve but leaves a *per-sequence* score shift averaging 3.01 MAD — and the average is the wrong statistic for a clinical claim: the 95th percentile is 8.65 MAD and the worst single held-out sequence moves 20.12, two fifths of the scale and 2.01 between-subject SD, purely because the camera moved. This is the difference between a guarantee and a fit. It is not an undertraining artefact: re-training the augmented arm at 3 $\times$ and 9 $\times$ the epoch budget leaves the shift flat (3.01/3.27/2.94 MAD mean at 60/180/540 epochs, worst sequence 20.1/21.4/19.9), so the residual swing is what augmentation converges to, not what it has yet to learn. The certificate is over Haar-random SO(3), so the 45° NTU RGB+D 60 Cross-View camera displacement is one instance of the group it covers: invariance there is *inherited* from Proposition 1, not separately measured. What NTU contributes independently is generalisation, not invariance — trained *without* rotation augmentation on cameras 2–3 and tested on the displaced camera 1, the identical encoder reaches 82.98% Top-1 against a faithful ST-GCN control’s 80.38%.

Clinical interpretation of the MAD scale. No inter-/intra-rater reliability coefficient (ICC, κ) is reported for KIMORE’s clinical Total Score [1] or for REHAB24-6’s; a targeted search, including direct attempts to retrieve the Capecci et al. full text (both the publisher record and the indexed copy returned access-blocked to automated retrieval), could not surface one, and we do not fabricate one — we flag its absence as this paper’s clearest open measurement gap. As external context, human inter-rater agreement on *movement-quality rating* tasks — the same construct KIMORE’s clinicians perform, scoring how well an exercise was executed rather than a functional-recovery outcome — typically falls in a moderate-to-good range: the Cutting Movement Assessment Score reports inter-rater ICC 0.58–0.91, and a cross-diagnostic Movement Quality Score reports ICC 0.93. We read this as a defensible band on what




Figure 2. MAD vs. camera azimuth. Certified-invariant models (EGRU, InvariantGRU, hand-crafted Ridge) are flat by theorem; all non-equivariant baselines (PCT, TCN, ST-GCN) degrade 9–13 MAD at extreme angles, indicating viewpoint fragility is a generic property of non-equivariant architectures. Augmentation flattens PCT’s aggregate but not its per-sequence swing (mean 3.01 MAD, worst sequence 20.12). Azimuth here is a rigid rotation of the ground-truth Kinect skeletons; it isolates the geometric effect of viewpoint and does not model per-view tracker degradation.

a KIMORE-specific coefficient would plausibly show, not a substitute for one. As a secondary, magnitude-of-change cross-check from a different construct and scale, minimal clinically important differences (MCIDs) on comparable 0–66/0–100-point motor-recovery scales sit in the low single digits to low teens of points (Fugl-Meyer upper-extremity 4–12.4, lower-extremity 6, Stroke Rehabilitation Assessment of Movement 1.9–4.8 per subscale), putting the augmented baseline’s per-sequence swing (mean 3.01, worst sequence 20.12, of 50) in the same order of magnitude as an established change-significance threshold on cognate instruments — at the tail, above it — and our own score shifts under rotation (9×10^{-6}), int8 and streaming (0.05 and lower, Sec. 5) orders of magnitude below it. A camera relocation that moves a patient’s score by more than an MCID is a measurement the instrument cannot be trusted to make; that is the practical content of the guarantee. The sharper, in-dataset question — does a 3.01-point swing cross KIMORE’s clinical bands? — has no answer: those bands do not exist, the five groups being recruitment categories, with a control scoring below a patient on every exercise (Es5: Expert 25.0, Parkinson 50.0). Measured against the score’s actual resolution, the mean swing is 0.30 between-subject SD and the worst sequence 2.01 SD, and it exceeds the gap separating 20.2% of subject pairs, 83.0% at the tail — pairs a camera can reorder (supplement). The MCID band then reads as corroboration from adjacent instruments, not as a KIMORE-specific threshold and not as the load-bearing

Table 4. Sensor-node failure: MAD as k nodes freeze (mean over 3 seeds $\times$ 5 folds; identical hash-locked inputs; no retraining). Underline: past the floor (8.31). “+oracle”: InvariantGRU given an oracle liveness mask (same checkpoint; features touching the dead joint zeroed at inference) — the measured fairness control.

Dead nodes k	EGRU (ours)	InvariantGRU	+oracle	PCT
1	7.10	<u>9.78</u>	<u>9.39</u>	6.76
4	8.41	<u>15.96</u>	<u>15.25</u>	7.51
8	<u>10.50</u>	<u>18.37</u>	<u>16.33</u>	8.73
MAD lost	+3.76	+12.05	+10.02	+2.27

claim.

Does the guarantee cost robustness? Sensor-node failure (Tab. 4). This is a tax check, not a selling point. Freezing k tracking nodes at their first-frame pose, the hand-crafted arm collapses through the floor at a *single* dead joint (6.31 $\rightarrow$ 9.78) because its features name their joints; EGRU pools learned graph messages and loses +3.76 MAD over 8 dead joints. **PCT wins this axis outright** (+2.27): it leads on KIMORE and, on REHAB24-6, overtakes us between $k=2$ and $k=4$ and finishes 0.021 AUROC ahead at $k=8$ (supplement). Node robustness is not a property of equivariance and we do not claim it. The table supports something narrower: the guarantee is not paid for with a robustness collapse ($k=8$ crosses the floor for every model, EGRU included). Both models see identically corrupted joints with no imputation, and we measured the fairness objection rather than assuming it: granting the hand-crafted features an *oracle* liveness mask (told exactly which joint died, no retraining) closes only $\sim 25\%$ of the gap to EGRU (+12.05 $\rightarrow$ +10.02 MAD lost; Tab. 4, right column) — the graph encoder needs no such signal and still loses far less, so the residual gap is architectural, not a missing-information artifact. The mechanism is invariant to the failure model (five operators incl. lag/burst/axis/lateral; supplement).

The mechanism is not an artifact of the freeze model. Under four further operators — stuck-at-lag, sporadic bursts, single-axis depth noise, and a left-limb-only variant — the ordering is invariant: the hand-crafted arm degrades most and the graph-pooled encoder least among the two exact-invariance models (Tab. 5). *Persistent* faults (freeze, lag) are far more destructive to naming features (+12.05, +8.88) than *transient* ones (burst +2.83, axis +1.93) — a length-masked temporal mean dilutes an intermittent spike but not a permanently corrupted coordinate — and unilateral (left-only) failure does not break the parity channels (InvariantGRU +8.73, EGRU +3.43).

The whole thesis in one grid (Tab. 6). *No model wins all five columns*, and that — not any single stress — is the claim. PCT is both the most accurate clean model (6.47) and the most node-robust (7.10); InvariantGRU is the cheapest at exactly zero degradation yet collapses at two

Table 5. Node failure under five operators (chiral models, 3 seeds $\times$ 5 folds, hash-locked). Each cell: MAD at $k=1$ / $k=8$ and total lost. The freeze/random row reproduces Tab. 4.

Failure operator	EGRU (ours)	InvariantGRU	PCT
freeze (random)	7.10/10.50 (+3.76)	9.78/18.37 (+12.05)	6.76/8.73 (+2.27)
freeze (left-only)	7.11/10.16 (+3.43)	10.81/15.04 (+8.73)	6.83/9.62 (+3.15)
stuck-at-lag	6.82/7.76 (+1.03)	9.02/15.19 (+8.88)	6.60/7.11 (+0.64)
sporadic burst	6.74/6.78 (+0.05)	6.69/9.14 (+2.83)	6.47/6.45 (−0.02)
axis/depth noise	6.77/7.20 (+0.47)	6.70/8.24 (+1.93)	6.48/6.50 (+0.03)

Table 6. Does each model clear the floor (8.31 MAD) under each stress? Last column: *mean* per-sequence score shift at the worst camera azimuth; the tail it conceals is in the text. EGNN shares the identical cut (viewpoint exact by construction). [†]Canon-PCA’s near-zero is *empirical*, a frame estimate stable on KIMORE, not a certificate.

Model	Params	Clean	90°	2 dead	Mean degr.
EGRU (ours)	0.66M	6.73	6.73✓	7.55✓	9×10^{-6}
InvariantGRU	0.21M	6.31	6.31✓	11.64×	0
Lighter $E(n)$ (EGNN)	~ 0.6 M	6.88	6.88✓	8.44×	1.4×10^{-5}
Canon-PCA + PCT	4.91M	6.85	6.85✓	7.79✓	$\approx 0^{\dagger}$
PCT (baseline)	4.91M	6.47	10.13×	7.10✓	9.33
PCT + rot	4.91M	7.47	7.68✓	7.76✓	3.01

dead nodes (11.64); and three models clear every stress — ours, Canon-PCA+PCT, and rotation-augmented PCT — so we are not the only survivor. Among those three, ours is the only one at 0.66M rather than 4.91M parameters and the only one whose invariance is a theorem rather than an estimated frame or an augmentation tolerance. What separates the survivors is the last column (Fig. 3a) — all flat in aggregate, but the augmented baseline’s mean per-sequence degradation is 3.01 MAD (worst sequence 20.12) against our 9×10^{-6} (3.0×10^{-4} worst), at $7.4\times$ more parameters. This is the only column where the difference is categorical rather than a trade: 3.01 is a tolerance that happens to be small on this data and has a 20 MAD tail, whereas 9×10^{-6} is roundoff around a value the theorem fixes at zero. Every other column is an axis we may lose, and on node failure we do lose it. Behind the *identical* invariant cut, a lighter $E(n)$ -equivariant (EGNN) encoder ties clean accuracy and viewpoint (exact by construction) but is measurably more node-fail brittle: tuned across depth, width, and coordinate-clamp over three seeds it still loses $+6.4 \pm 1.3$ MAD over eight dead joints (vs. EGRU’s +3.76) and crosses the floor already at two dead joints (8.44 vs. EGRU’s 7.55) — so the steerable encoder is the more robust of the two equivariant options here.

Canonicalization ties us offline — and the difference is the guarantee. The strongest objection to an equivariant model is that a non-equivariant backbone fed a canonical frame reaches the same place for less. It does, here: a per-frame PCA principal-axis frame feeding the transformer (Canon-PCA+PCT, Tab. 6) matches every offline axis — clean 6.85 MAD, 7.79 at two dead joints, and a viewpoint invariance empirically exact for arbitrary $SO(3)$,




Figure 3. **The axis that separates the survivors.** (a) Per-patient score shift under camera rotation (log scale). Bar is the mean over held-out sequences at the worst azimuth; whisker runs 95th percentile to worst sequence. EGNN has no measured tail and carries no whisker. (b) PCA frame instability: frame transitions flipping > 45° (21%), and frames within an eigen-gap of 0.05 of degeneracy (18%).

not merely the azimuth subgroup a clinician sweeps (worst canonical-coordinate shift 3.7×10^{-11} over 32 random rotations). This is the strong form of the baseline: its axes are sign-disambiguated by a deterministic rule (each oriented so its dominant joint projection is positive) and the frame forced right-handed, so the discrete $\pm$ ambiguity is resolved before any number below is measured.

The difference is not accuracy but *conditioning*. Under a clean rotation the PCA frame is exactly equivariant — rotating the input sends the covariance to RCR^T , whose eigenvectors are exactly RV — and we measure that residual at 3×10^{-6} degrees. But by Davis–Kahan [16] an estimated eigenvector subspace moves by $O(\|E\|/\gamma)$ under a perturbation E with eigen-gap γ , and a depth sensor supplies $E \neq 0$ at every frame. On KIMORE γ is small far more often than not — 18% of frames within 0.05 of degeneracy and 45% within 0.10, a standing body being nearly axially symmetric — so amplification is the common case, not a corner. Injecting camera-frame noise at 2% of body scale and comparing the frame recovered at two camera poses, the residual rotation grows from 1.2° where $\gamma > 0.25$ to 11.1° where $\gamma < 0.02$: a $9\times$ amplification, with 14.5% of frames past 45° .

A better sign rule does not fix it. Replacing the dominant-joint convention with a third-moment rule leaves the amplification intact (12.1° , trimming only catastrophic flips to 11.1°), because a sign rule resolves a *discrete* ambiguity whereas degeneracy is a *continuous* one. Temporal smoothing is worse (28.8° ; 31.6% past 45°): a frame locked onto the wrong branch persists, and the frame becomes

history-dependent, trading per-frame equivariance for lag. A learned point-estimate frame inherits the same obstruction — no canonicalizer is simultaneously exactly equivariant and continuous where the configuration’s stabilizer is non-trivial [17]. We deliberately report this at the level of the frame rather than the score: the frame measurement is model-free and therefore binds any spectral canonicalizer, whereas how much of the wobble survives to the output depends on the corpus and on how much jitter the downstream network has learned to absorb — a dataset-dependent question, and secondary to whether the mechanism exists. Our cut estimates *no* frame: it is $SE(3)$ -invariant for arbitrary weights by construction (Proposition 1), certified to round-off over the full rotation group — so the *same* certificate covers, with no retraining, the genuinely displaced 45° camera of NTU RGB+D Cross-View. A frame estimator offers neither the worst-case bound nor that transfer. We do not claim an accuracy or on-distribution win over the hand-crafted or canonicalized baselines here — the tie is total; we claim a *certified* worst-case that they structurally cannot offer, and we recommend the cheaper models where an off-distribution guarantee is not required.

When the worst case is the deployment case. Concretely: a home tele-rehabilitation device is set up once by the patient, not a technician, and left running unattended — the camera sits wherever a phone stand or a laptop lid happens to land, and the consumer depth sensor is the component most likely to lose a joint (a hand behind the back, a node that freezes mid-repetition). Neither cheaper alternative is safe against *both* conditions at once. Canon-PCA+PCT’s viewpoint guarantee is only empirical: Fig. 3b shows the PCA frame sits within an eigen-gap of 0.05 of degeneracy on 18% of frames — a fast or asymmetric repetition, a partially occluded limb — and a flipped frame there is silent, no flag, just a wrong score. InvariantGRU’s viewpoint guarantee *is* exact, like ours, but its named-joint features make it the most node-fail-fragile model measured: one frozen tracking node drives it from 6.31 to 9.78 MAD, through the floor (Tab. 4), the same freeze mode a consumer sensor exhibits routinely. Uncontrolled placement and a tracking node that drops are exactly the two conditions a home device cannot rule out in advance, and EGRU is the only model measured that is provably safe against both at once. That is the niche this paper stakes out: not that InvariantGRU or Canon-PCA+PCT are worse choices in a supervised clinic with a fixed rig and a technician present, but that neither offers a bound that survives an unsupervised deployment on its own.

Second-corpus replication (REHAB24-6). To de-risk single-corpus reliance we replicate the structural story on REHAB24-6 (binary correct/incorrect repetitions; 1057 reps, 10 subjects; the identical 0.66M model as a BCE logit), an AUROC task over 3 seeds and subject-disjoint CV

Table 7. REHAB24-6 second-corpus replication (3 seeds, subject-disjoint CV). AUROC@90° is the same model evaluated under a 90° camera rotation: EGRU is invariant by theorem (logit drift $\leq 1.5 \times 10^{-4}$); the transformer falls to chance (drift 23.2). Node-fail: AUROC lost over 8 frozen nodes. The two models *cross* between 4 and 8 frozen joints; full curve and fold-level spread in the supplement.

Model	AUROC	AUROC@90°	Node-fail lost
EGRU (ours)	0.738 ± 0.01	0.738	+0.14
PCT (baseline)	0.707 ± 0.02	0.52×	+0.08

(Tab. 7). The accuracy tie holds — EGRU 0.738 ± 0.01 versus the transformer 0.707 ± 0.02 — and, decisively, the viewpoint theorem transfers: under a 90° camera rotation EGRU’s AUROC is *unchanged* (logit drift $\leq 1.5 \times 10^{-4}$), while the transformer collapses to chance (0.52). Node failure degrades both gracefully. The guarantee is thus not a KIMORE artifact — though at 10 subjects this corpus replicates *structural* properties, not accuracy: a certified invariance either transfers exactly or it does not, and we draw no accuracy conclusion from it.

Per-family contribution of the cut. Zeroing one family of Eq. (3) at inference on the deployed chiral model (dimension preserved, no retraining) resolves the mechanism to family granularity (full table in the supplement). Two families carry accuracy — the pairwise *cosines* (+1.83 MAD to remove) and the learned *triple products* (+2.75) — and removing the cosines is doubly harmful, inflating node-failure degradation the most (+7.29). The parity-odd families *split*: the learned triple products are relied upon while the fixed anatomical volumes are inert (+0.19), which reconciles rather than contradicts the chirality null — that null is a *retrain* comparison (parity-even families substitute during training), whereas this is an *inference* ablation on a model trained *with* the family.

Two negative results, honestly — each a mapped boundary of where structure pays. (i) Native consumption of irregular arrivals does not help on KIMORE — the fixed-grid baseline is better at every drop level under a hash-locked budget — and a spectral census [23] predicts it: 97.7% of *positional* energy lies below the resampling corner (though 70% of *velocity* energy lies above it), so resampling preserves pose geometry while low-passing the derivative band — on KIMORE’s healthy-form exercises, where the discriminative signal is positional, this acts as a denoiser; the boundary would move for tremor-dominated tasks. The mechanism is real (a 5 Hz tremor is recovered at $r=1.000$ through 70% drops where a resampled estimator decays to 0.785), so the claim is a scope boundary, not a defeat. (ii) Restoring chirality is principled but not rewarded on healthy-form exercises with no chiral pathology — the boundary at which a handedness prior stops paying. Both are stated with the corpus properties that would reward them

Table 8. Deployment audit. Top: the *full* equivariance precision budget of the deployed pooled EGRU (INV floor = relative violation of the 160-d parity-even projection; score shift in MAD of 50). fp16 is 3.2× tighter than bf16 — equivariance pays in mantissa bits, which bf16 trades for exponent range — and the int8 cliff is *solely* activation grid-rounding (weight-only quant preserves the theorem). Bottom: causal streaming vs. bidirectional read-out.

Arithmetic	INV floor	Score shift	Note
fp64	5.75×10^{-16}	3.3×10^{-14}	exact (theorem)
fp32	2.68×10^{-7}	8.9×10^{-6}	= main viewpoint
fp16	2.62×10^{-3}	0.024	> bf16
bf16	8.34×10^{-3}	0.195	mantissa cost
int8, weights only	3.09×10^{-7}	1.2×10^{-5}	theorem holds
int8, activations only	3.38×10^{-2}	0.056	cliff (acts)
int8, weights+acts	3.40×10^{-2}	0.051	weight-quant. theorem-safe

Streaming: 8.9 ms/frame (112 fps); read-out cut unchanged;
causal vs. bidir NTU X-View 79.8 vs. 83.0 (−3.2 pt).

(supplement).

5. Deployment Analysis

We treat the guarantee as a systems property (Tab. 8). *Precision budget*: the invariance floor climbs monotonically from fp64 to int8, and the int8 cliff is attributable *solely* to activation grid-rounding of the steerable channels — weight quantization preserves the theorem exactly (a quantized model is a different equivariant model) — while the reported clinical score still moves only 0.05 MAD — below the floor, 0.005 of the between-subject SD, and one to two orders of magnitude below the comparator MCID band established in Sec. 4. *Streaming*: because invariance lives in a per-frame encoder and an invariant cut, the read-out converts to a strictly causal unidirectional form with a running-mean pool that emits a first score 8.9 ms after the first frame on an RTX-class GPU (batch 1; 112 fps, inside a 30 fps budget) — we report GPU latency as an upper bound on compute and leave edge-SoC profiling to deployment — reaching a usable prediction after ~ 55 frames (median), an honest time-to-first-*reliable*-score for real-time use. The invariant read-out cut is unchanged by the causal pooling, so Proposition 1 carries over verbatim; the cost is 3.2 points of NTU accuracy vs. the bidirectional model (NTU-60 X-View SOTA is $\sim 92\%$; our backbone is deliberately untuned — we measure the streaming *delta*, not absolute accuracy). A canonicalization front-end would instead add a per-frame eigendecomposition to the hot path and re-open the invariance question after quantization; here the guarantee is a fixed property of the read-out and survives both int8 and causal streaming without re-certification.

6. Conclusion

On the metric the field reports, no architecture separates from another: on clean frontal KIMORE our model, the baseline, and a hand-crafted-invariant GRU tie within a

0.33 MAD nondeterminism floor — the metric is saturated. What separates them is off the frontal manifold. Our read-out is $SE(3)$ -invariant by construction and certified to machine roundoff, so relocating the camera moves a patient’s score by 9×10^{-6} MAD for arbitrary weights, while a broken tracking node — fatal to the hand-crafted model — is absorbed by the graph encoder; both properties replicate on a second corpus, and the guarantee survives to int8 and to causal streaming. We recommend the field report a mean-predictor floor and the per-sequence degradation *distribution* — not only its mean, which on our own augmented baseline understates the worst patient by nearly $7 \times$ — alongside aggregate error, and evaluate off the frontal manifold: camera relocation and node failure are ordinary conditions of a home device, not adversarial constructions. **Code availability.** KIMORE [1], NTU RGB+D and REHAB24-6 are public. Our certificate, harnesses, benchmarks, and the script regenerating every number here from banked artifacts will be released on acceptance; the URL is withheld for double-blind review.

References

- [1] M. Capecci et al., “The KIMORE dataset: Kinematic assessment of movement and clinical scores for remote monitoring of physical rehabilitation,” *IEEE TNSRE*, vol. 27, no. 7, pp. 1436–1448, 2019.
- [2] K. Rafat, M. I. Hossain, M. M. L. Elahi, S. Momen, F. Rahman, N. Mohammed, and S. Rahman, “A point cloud transformer for remote monitoring and automated assessment of physical rehabilitation exercises,” *IEEE J. Biomed. Health Inform.*, 2026, accepted; arXiv:2606.30309.
- [3] M. Karlov, A. Abedi, and S. S. Khan, “Rehabilitation exercise quality assessment through supervised contrastive learning with hard and soft negatives,” *Med. Biol. Eng. Comput.*, 2024.
- [4] A. Abedi, M. Malmirian, and S. S. Khan, “Cross-modal video to body-joints augmentation for rehabilitation exercise quality assessment,” arXiv:2306.09546, 2023.
- [5] Z. Kuang et al., “Dual-stream STGCN with motion-aware grouping for rehabilitation action quality assessment,” *Sensors*, vol. 26, no. 1, p. 287, 2026.
- [6] A. Ismail-Fawaz et al., “Rehab-Pile: A cross-subject rehabilitation benchmark aggregating KIMORE, UI-PRMD and IRDS,” arXiv:2507.21018, IEEE FG 2026.
- [7] A. Vakanski et al., “A data set of human body movements for physical rehabilitation exercises (UI-PRMD),” *Data*, vol. 3, no. 1, p. 2, 2018.
- [8] S. Yan, Y. Xiong, and D. Lin, “Spatial temporal graph convolutional networks for skeleton-based action recognition,” in *Proc. AAAI*, 2018.
- [9] T. S. Cohen and M. Welling, “Group equivariant convolutional networks,” in *Proc. ICML*, 2016.
- [10] N. Thomas et al., “Tensor field networks: Rotation- and translation-equivariant neural networks for 3D point clouds,” arXiv:1802.08219, 2018.
- [11] F. B. Fuchs et al., “SE(3)-Transformers: 3D rotation-translation equivariant attention networks,” in *Proc. NeurIPS*, 2020.
- [12] M. Geiger and T. Smidt, “e3nn: Euclidean neural networks,” arXiv:2207.09453, 2022.
- [13] V. G. Satorras, E. Hoogeboom, and M. Welling, “E(n) equivariant graph neural networks,” in *Proc. ICML*, 2021.
- [14] C. Deng et al., “Vector neurons: A general framework for SO(3)-equivariant networks,” in *Proc. ICCV*, 2021.
- [15] O. Puny et al., “Frame averaging for invariant and equivariant network design,” in *Proc. ICLR*, 2022.
- [16] C. Davis and W. M. Kahan, “The rotation of eigenvectors by a perturbation. III,” *SIAM J. Numer. Anal.*, vol. 7, no. 1, pp. 1–46, 1970.
- [17] S.-O. Kaba et al., “Equivariance with learned canonicalization functions,” in *Proc. ICML*, 2023.
- [18] R. T. Q. Chen et al., “Neural ordinary differential equations,” in *Proc. NeurIPS*, 2018.
- [19] Y. Rubanova, R. T. Q. Chen, and D. Duvenaud, “Latent ODEs for irregularly-sampled time series,” in *Proc. NeurIPS*, 2019.
- [20] P. Kidger et al., “Neural controlled differential equations for irregular time series,” in *Proc. NeurIPS*, 2020.
- [21] Z. Che et al., “Recurrent neural networks for multivariate time series with missing values,” *Sci. Rep.*, vol. 8, p. 6085, 2018.
- [22] J. R. Dormand and P. J. Prince, “A family of embedded Runge–Kutta formulae,” *J. Comput. Appl. Math.*, vol. 6, no. 1, pp. 19–26, 1980.
- [23] J. T. VanderPlas, “Understanding the Lomb–Scargle periodogram,” *ApJS*, vol. 236, no. 1, p. 16, 2018.